

CPSC 340: Machine Learning and Data Mining

Generative Models

Admin

- **Assignment 0** was due last Wednesday.
 - Because of late days, we have to wait 3 days to post solutions.
 - At that time you will also gain read access to your classmates' work.
 - We voted on this during the first lecture.
 - No one approached me privately with objections.
- **Assignment 1** is out.
 - This is a representative assignment w.r.t. length/difficulty/format/style.
- **Registration:**
 - Keep checking your registration, it could change quickly.
 - As of last night, waitlist was down to 14 people.
- **Probability:**
 - If you are struggling with probability concepts towards the end of class today, check out the posted notes on probability.

Last Time: Training, Testing, and Validation

- Training step:

Input: set of ' n ' training examples x_i with labels y_i

Output: a model that maps from arbitrary x_i to a y_i

- Prediction step:

Input: set of ' t ' testing examples \hat{x}_i and a model.

Output: predictions \hat{y}_i for the testing examples.

- What we are interested in is the **test error**:

- Error made by prediction step on new data.

- Validation set or cross-validation can be used to estimate test error.

Should you trust them?

- Scenario 1:
 - “I built a model based on the data you gave me.”
 - “It classified your data with 98% accuracy.”
 - “It should get 98% accuracy on the rest of your data.”
- Probably not:
 - They are reporting training error.
 - This might have nothing to do with test error.
 - E.g., they could have fit a very deep decision tree.
- Why ‘probably’?
 - If they only tried a few very simple models, the 98% might be reliable.
 - E.g., they only considered decision stumps with simple 1-variable rules.

Should you trust them?

- Scenario 2:
 - “I built a model based on half of the data you gave me.”
 - “It classified the other half of the data with 98% accuracy.”
 - “It should get 98% accuracy on the rest of your data.”
- Probably:
 - They computed the validation error once.
 - This is an unbiased approximation of the test error.
 - Trust them if you believe they didn’t violate the golden rule.

Should you trust them?

- Scenario 3:
 - “I built 10 models based on half of the data you gave me.”
 - “One of them classified the other half of the data with 98% accuracy.”
 - “It should get 98% accuracy on the rest of your data.”
- Probably:
 - They computed the validation error a small number of times.
 - Maximizing over these errors is a biased approximation of test error.
 - But they only maximized it over 10 models, so bias is probably small.
 - They probably know about the golden rule.

Should you trust them?

- Scenario 4:
 - “I built 1 billion models based on half of the data you gave me.”
 - “One of them classified the other half of the data with 98% accuracy.”
 - “It should get 98% accuracy on the rest of your data.”
- Probably not:
 - They computed the validation error a huge number of times.
 - Maximizing over these errors is a biased approximation of test error.
 - They tried so many models, one of them is likely to work by chance.
 - This is the “multiple comparisons problem” in statistics
- Why ‘probably’?
 - If the 1 billion models were all extremely-simple, 98% might be reliable.

Should you trust them?

- Scenario 5:
 - “I built 1 billion models based on the first third of the data you gave me.”
 - “One of them classified the second third of the data with 98% accuracy.”
 - “It also classified the last third of the data with 98% accuracy.”
 - “It should get 98% accuracy on the rest of your data.”
- Probably:
 - They computed the first validation error a huge number of times.
 - But they had a second validation set that they only looked at once.
 - The second validation set gives unbiased test error approximation.
 - This is ideal, as long as they didn't violate golden rule on second set.
 - And assuming you are using IID data in the first place.

The 'Best' Machine Learning Model

- Decision trees are not always most accurate.
- What is the 'best' machine learning model?
- First we need to define generalization error:
 - Test error on new examples (excludes test examples seen during training).
- No free lunch theorem:
 - There is **no** 'best' model achieving the best generalization error for every problem.
 - If model A generalizes better to new data than model B on one dataset, there is another dataset where model B works better.
- This question is like asking which is 'best' among "rock", "paper", and "scissors".

The 'Best' Machine Learning Model

- Implications of the lack of a 'best' model:
 - We need to learn about and **try out multiple models**.
- So which ones to study in CPSC 340?
 - We'll usually motivate a method by a specific application.
 - But we'll focus on **models that are effective in many applications**.
- Caveat of no free lunch (NFL) theorem:
 - The world is very structured.
 - **Some datasets are more likely than others**.
 - Model A really could be better than model B on every real dataset in practice.
- Machine learning research:
 - Large focus on models that are **useful across many applications**.

Application: E-mail Spam Filtering

- Want a build a system that filters spam e-mails.

<input type="checkbox"/>			Jannie Keenan	ualberta	You are owed \$24,718.11
<input type="checkbox"/>			Abby	ualberta	USB Drives with your Logo
<input type="checkbox"/>			Rosemarie Page		Re: New request created with ID: ##62
<input type="checkbox"/>			Shawna Bulger		RE: New request created with ID: ##63
<input type="checkbox"/>			Gary	ualberta	Cooperation



Gary <jaiwasie@mail.com>

to schmidt ▾



Be careful with this message. Similar messages were used to steal people's personal information. [Learn more](#)

Hey,

Do you have a minute today?

Are you interested to use our email marketing and lead generation solutions?

We have worked on a number of projects and campaigns in many industries since 2007

Please reply today so we can go over options for you.

I am sure we can help to grow your business soon by using our mailing services.

Best regards,

Gary

Contact: abelfong@sina.com

- We have a big collection of e-mails, labeled by users.
- Can we formulate as supervised learning?

First a bit more supervised learning notation

- We have been using the notation 'X' and 'y' for supervised learning:

The diagram illustrates the notation for supervised learning. It shows a matrix X and a vector y . A red circle highlights a row in X , with an arrow pointing to x_i , representing the features of object i . Another red circle highlights an element in y , with an arrow pointing to y_i , representing the label of object i . A red arrow points from the row in X to the element in y , indicating the mapping from features to labels. A red arrow also points from the label x_{ij} to a specific element in the row of X , representing feature j of object i .

- X is matrix of all features, y is vector of all labels.
- Need a way to refer to the features and label of **specific object 'i'**.
 - We use y_i for the label of object 'i' (element 'i' of 'y').
 - We use x_i for the features object 'i' (row 'i' of 'X').
 - We use x_{ij} for feature 'j' of object 'i'.

Feature Representation for Spam

- How do we make label ' y_i ' of an individual e-mail?
 - ($y_i = 1$) means 'spam', ($y_i = 0$) means 'not spam'.
- How do we construct features ' x_i ' for an e-mail?
 - Use **bag of words**:
 - "hello", "vicodin", "\$".
 - "vicodin" feature is 1 if "vicodin" is in the message, and 0 otherwise.
 - Could add phrases:
 - "be your own boss", "you're a winner", "CPSC 340".
 - Could add regular expressions:
 - <recipient>, <sender domain == "mail.com">

Probabilistic Classifiers

- For years, best spam filtering methods used **naïve Bayes**.
 - Naïve Bayes is a **probabilistic** classifier based on **Bayes rule**.
 - It's "naïve" because it makes a **strong conditional independence assumption**.
 - But it tends **to work well with bag of words**.
- Probabilistic classifiers model the **conditional probability**, $p(y_i | x_i)$.
 - "If a message has words x_i , what is probability that message is spam?"
- If $p(y_i = \text{'spam'} | x_i) > p(y_i = \text{'not spam'} | x_i)$, classify as spam.

- Recall our spam filtering setup:
 - y_i : whether or not the e-mail was spam.
 - x_i : the **set of words/phrases/expressions** in the e-mail.
- To model conditional probability, **naïve Bayes** uses Bayes rule:

$$p(y_i = \text{"spam"} \mid x_i) = \frac{p(x_i \mid y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- Easy part #1: $p(y_i = \text{'spam'})$ is the **probability that an e-mail is spam**.
 - Count of number of times ($y_i = \text{'spam'}$) divided by number of objects 'n'.

- Recall our spam filtering setup:
 - y_i : whether or not the e-mail was spam.
 - x_i : the **set of words/phrases/expressions** in the e-mail.
- To model conditional probability, **naïve Bayes** uses Bayes rule:

$$p(y_i = \text{"spam"} | x_i) = \frac{p(x_i | y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- Easy part #2: We **don't need** $p(x_i)$.

To test $p(y_i = \text{"spam"} | x_i)$ we just need to know if $p(y_i = \text{"spam"} | x_i) > p(y_i = \text{"not spam"} | x_i)$.

By Bayes rule this is equivalent to $\frac{p(x_i | y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)} > \frac{p(x_i | y_i = \text{"not spam"}) p(y_i = \text{"not spam"})}{p(x_i)}$

Denominators are the same so we just test $p(x_i | y_i = \text{"spam"}) p(y_i = \text{"spam"}) > p(x_i | y_i = \text{"not spam"}) p(y_i = \text{"not spam"})$

Generative Classifiers

- The **hard part is estimating $p(x_i \mid y_i = \text{'spam'})$** :
 - the probability of seeing the words/expressions x_i if the e-mail is spam.
- Classifiers based on Bayes rule are called **generative classifier**:
 - It needs to know the **probability of the features, given the class**.
 - How to “generate” features.
 - You need a model that knows what spam messages look like.
 - And a second that knows what non-spam messages look like.
 - This work well with **tons of features compared to number of objects**.

Generative Models

- Spam filtering methods based on **generative models**:

$$p(y_i = \text{"spam"} | x_i) = \frac{p(x_i | y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- What do these terms mean?

ALL E-MAILS
(including duplicates)

Generative Models

- Spam filtering methods based on **generative models**:

$$p(y_i = \text{"spam"} \mid x_i) = \frac{p(x_i \mid y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- $p(x_i)$ is probability that a random e-mail has features x_i .

ALL E-MAILS
(including duplicates)

Generative Models

- Spam filtering methods based on **generative models**:

$$p(y_i = \text{"spam"} \mid x_i) = \frac{p(x_i \mid y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- $p(x_i)$ is probability that a random e-mail has features x_i .

ALL E-MAILS
(including duplicates)

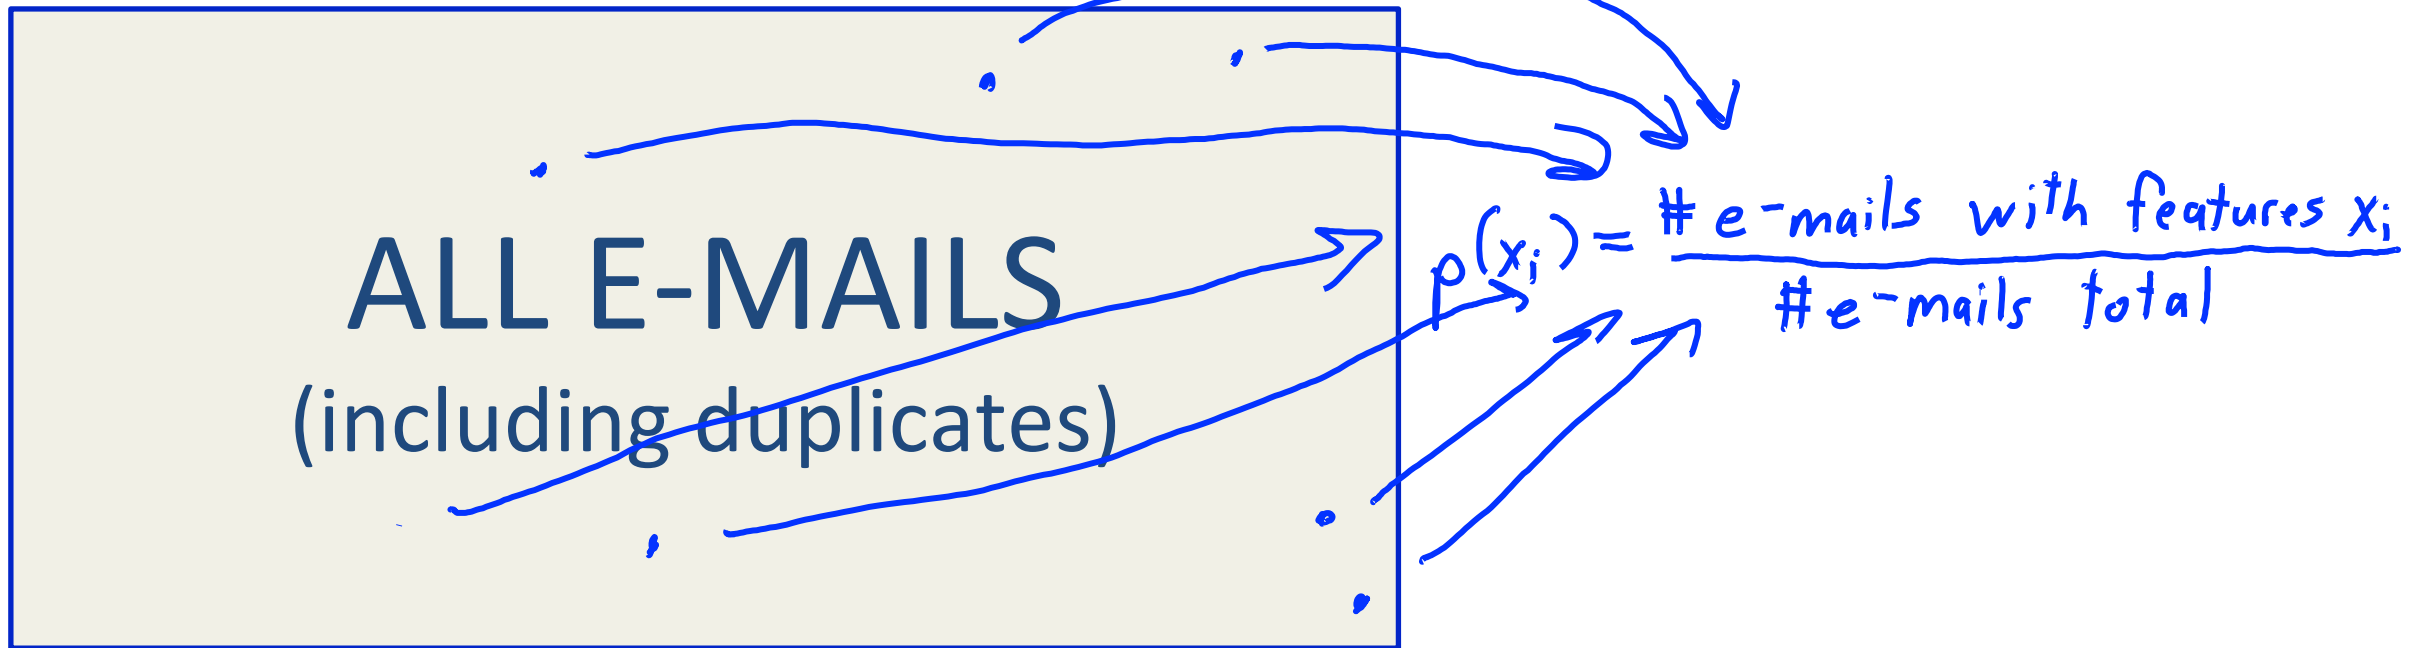
$$p(x_i) = \frac{\# \text{ e-mails with features } x_i}{\# \text{ e-mails total}}$$

Generative Models

- Spam filtering methods based on **generative models**:

$$p(y_i = \text{"spam"} | x_i) = \frac{p(x_i | y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- $p(x_i)$ is probability that a random e-mail has features x_i .



Generative Models

- Spam filtering methods based on **generative models**:

$$p(y_i = \text{"spam"} \mid x_i) = \frac{p(x_i \mid y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- $p(x_i)$ is probability that a random e-mail has features x_i .

ALL E-MAILS
(including duplicates)

$$p(x_i) = \frac{\# \text{e-mails with features } x_i}{\# \text{e-mails total}}$$

- Hard, but not needed to classify using:
 $p(y_i = \text{'spam'} \mid x_i) > p(y_i = \text{'not spam'} \mid x_i)$

Generative Models

- Spam filtering methods based on **generative models**:

$$p(y_i = \text{"spam"} | x_i) = \frac{p(x_i | y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- $p(y_i = \text{'spam'})$ is probability that a random e-mail is spam.



$$p(y_i = \text{"spam"}) = \frac{\# \text{ spam messages}}{\# \text{ total messages}}$$

- Hard to compute exactly.
- But is easy to approximate from data:
 - Count (#spam in data)/(#messages)

Generative Models

- Spam filtering methods based on **generative models**:

$$p(y_i = \text{"spam"} | x_i) = \frac{p(x_i | y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- $p(x_i | y_i = \text{'spam'})$ is probability that spam has features x_i .




$$p(x_i | y_i = \text{"spam"}) = \frac{\# \text{ spam messages with features } x_i}{\# \text{ spam messages}}$$

Generative Models

- Spam filtering methods based on **generative models**:

$$p(y_i = \text{"spam"} | x_i) = \frac{p(x_i | y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- $p(x_i | y_i = \text{'spam'})$ is probability that spam has features x_i .



- Very hard to estimate:**
 - Too many possible x_i .

Naïve Bayes

- How the naïve Bayes model deals with the hard terms:

$$p(\text{spam} | \text{hello}, \text{vicodin}, \text{CPSC 340}) = \frac{p(\text{hello}, \text{vicodin}, \text{CPSC 340} | \text{spam}) p(\text{spam})}{p(\text{hello}, \text{vicodin}, \text{CPSC 340})} \quad (\text{Bayes rule})$$


"equal up to constant not depending on spam" \propto $\underbrace{p(\text{hello}, \text{vicodin}, \text{CPSC 340} | \text{spam})}_{\text{HARD}} \underbrace{p(\text{spam})}_{\text{easy}}$

(naive Bayes assumption) $\approx \underbrace{p(\text{hello} | \text{spam})}_{\text{easy}} \underbrace{p(\text{vicodin} | \text{spam})}_{\text{easy}} \underbrace{p(\text{CPSC 340} | \text{spam})}_{\text{easy}} \underbrace{p(\text{spam})}_{\text{easy}}$

- Now **only** need easy quantities like $p(\text{'vicodin'} = 1 | y_i = \text{'spam'})$.

Naïve Bayes Models

- $p(\text{vicodin} = 1 \mid \text{spam} = 1)$ is probability of seeing 'vicodin' in spam.



$$p(\text{vicodin}=1 \mid \text{spam}=1) = \frac{\# \text{ spam messages w/ vicodin}}{\# \text{ spam messages}}$$

- Easy to estimate:
 - $\#(\text{spam w/ Vicodin}) / \# \text{spam}$
 - "Maximum likelihood estimate"

- In naïve Bayes: assume features are independent given label.
 - "Once you know it's spam, there is no dependency between features."
 - Not true, but sometimes a good approximation.

Naïve Bayes

- Naïve Bayes more formally:

$$p(y_i | x_i) = \frac{p(x_i | y_i) p(y_i)}{p(x_i)}$$

$$\propto p(x_i | y_i) p(y_i)$$

$$\approx \prod_{j=1}^d [p(x_{ij} | y_i)] p(y_i)$$

- Assumption: all x_i are **conditionally independent** given y_i .

Independence of Random Variables

- **Events** A and B are **independent** if $p(A, B) = p(A)p(B)$.
 - Equivalently: $p(A | B) = p(A)$.
 - “Knowing B happened tells you nothing about A”.
 - We use the notation:
$$A \perp B$$
- **Random variables** are **independent** if $p(x, y) = p(x)p(y)$ for all x and y.
 - Flipping two coins:
$$p(C_1 = \text{'heads'}, C_2 = \text{'heads'}) = p(C_1 = \text{'heads'})p(C_2 = \text{'heads'})$$
$$p(C_1 = \text{'tails'}, C_2 = \text{'heads'}) = p(C_1 = \text{'tails'})p(C_2 = \text{'heads'})$$
$$\dots$$

Conditional Independence

- Example: food poisoning
 - If food was bad, each person independently gets sick with probability 50%
 - Unconditionally, me getting and and you getting sick are NOT independent
 - If I got sick, that makes me think the food was bad, which makes it more likely that you will get sick also. So knowing my situation influences my beliefs about yours.
 - But, conditioned on knowing the food was bad (or not bad), my sickness and your sickness are independent.
- Definition: A and B are **conditionally independent** given C if
$$p(A, B \mid C) = p(A \mid C)p(B \mid C).$$
 - Equivalently: $p(A \mid B, C) = p(A \mid C)$.
 - “Knowing C happened, also knowing B happened says nothing about A”.
 - We use the notation: $A \perp B \mid C$

Naïve Bayes for any number of classes

- Let c be a class label in $\{c_1, c_2, \dots\}$
 - In the spam example, we only had 2 classes (spam and not spam)
- Let i be a training example's index, j a feature index, k a feature value

Training:

1. Set n_c to the number of times $(y = c)$.
2. Estimate $p(y = c)$ as $\frac{n_c}{n}$.
3. Set n_{cjk} as the number of times $(y_i = c, X_{ij} = k)$
4. Estimate $p(x_i = k \mid y = c) = \frac{p(x_i = k, y = c)}{p(y = c)}$ as $\frac{\frac{n_{cjk}}{n}}{\frac{n_c}{n}} = \frac{n_{cjk}}{n_c}$.

Naïve Bayes for any number of classes

Prediction:

Given a new example x_i we want to find the 'c' maximizing $p(x_i | y_i)$.

Under the naive Bayes assumption we thus maximize

$$p(y=c | x_i) \propto \prod_{j=1}^d [p(x_{ij} | y=c)] p(y=c)$$

- Note that these terms do not add up to 1 because we dropped the denominator $p(x_i)$.

Application: E-mail Spam Filtering

- Want to build a system that filters spam e-mails:
- We formulated as supervised learning:
 - $(y_i = 1)$ if e-mail 'i' is spam, $(y_i = 0)$ if e-mail is not spam.
 - $(x_{ij} = 1)$ if word/phrase 'j' is in e-mail 'i', $(x_{ij} = 0)$ if it is not.

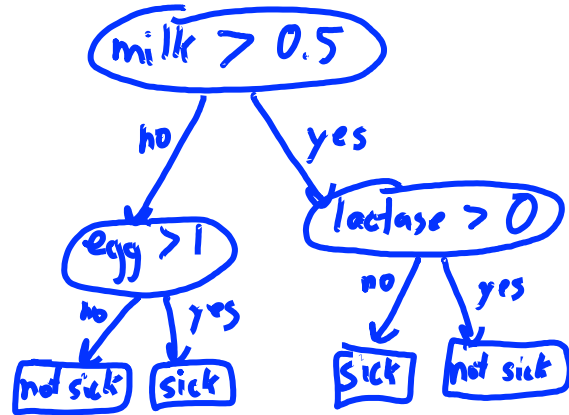
<input type="checkbox"/>			Jannie Keenan	ualberta	You are owed \$24,718.11
<input type="checkbox"/>			Abby	ualberta	USB Drives with your Logo
<input type="checkbox"/>			Rosemarie Page		Re: New request created with ID: ##62
<input type="checkbox"/>			Shawna Bulger		RE: New request created with ID: ##63
<input type="checkbox"/>			Gary	ualberta	Cooperation

\$	Hi	CPSC	340	Vicodin	Offer	...
1	1	0	0	1	0	...
0	0	0	0	1	1	...
0	1	1	1	0	0	...
...

Spam?
1
1
0
...

Decision Trees vs. Naïve Bayes

- Decision trees:



1. Sequence of rules based on 1 feature.
2. Training: 1 pass over data per depth.
3. Hard to find optimal tree.
4. Testing: just look at features in rules.
5. New data: might need to change tree.
6. Accuracy: good if simple rules work.

- Naïve Bayes:

$$p(\text{sick} \mid \text{milk}, \text{egg}, \text{lactase}) \\ \approx p(\text{milk} \mid \text{sick}) p(\text{egg} \mid \text{sick}) p(\text{lactase} \mid \text{sick}) p(\text{sick})$$

1. Simultaneously combine all features.
2. Training: 1 pass over data to count.
3. Easy to find optimal probabilities.
4. Testing: look at all features.
5. New data: just update counts.
6. Accuracy: good if features almost independent given label.

Naïve Bayes Issues

1. Do we **need to store the full bag of words** 0/1 variables?
 - No: only need **list of non-zero features** for each e-mail.
 - Could use a **sparse matrix** representation.
2. Problem with maximum likelihood estimate (MLE):
 - MLE of $p(\text{'lactase'} = 1 \mid \text{'spam'})$ is $(\text{\#spam messages with 'lactase'})/\text{\#spam}$.
 - If you have **no spam messages with lactase**:
 - $p(\text{'lactase'} \mid \text{'spam'}) = 0$, and message automatically gets through filter.
 - Fix: imagine we saw/not-saw each word in spam/not-spam messages:
 - “Laplace smoothing”: assume some “pseudo-counts” for each feature/label.
 - for binary features: replace n_{cjk}/n_c with $(n_{cjk} + 1)/(n_c + 2)$.
 - a generalization is $(n_{cjk} + \beta)/(n_c + 2\beta)$ for some constant β .
 - If X_{ij} can take ‘m’ values, you would $(n_{cjk} + \beta)/(n_c + m\beta)$

Naïve Bayes Issues

3. During the prediction, the **probability can underflow**:

$$p(y=c | x_i) \propto \prod_{j=1}^d [p(x_{ij} | y=c)] p(y=c)$$

→ All these are < 1 so the product gets very small!

- Standard fix is to (equivalently) **maximize the logarithm of the probability**:
 - Logarithm **turns multiplication of small numbers into addition** of small numbers.
 - Logarithm is monotonic, so it **doesn't change location of the maximum (maximizer)**
 - See CPSC 302/303 for more on underflow and floating point issues.

Decision Theory

- Spam classification example
 - Are we **equally concerned about spam vs. not spam?**
- True positives, false positives, false negatives, false negatives:

Predict / True	True 'spam'	True 'not spam'
Predict 'spam'	True Positive	False Positive
Predict 'not spam'	False Negative	True Negative

- The costs mistakes might be different:
 - Letting a spam message through (false negative) is not a big deal.
 - Filtering a not spam (false positive) message will make users mad.

Decision Theory

- We can give a **cost** to each scenario, such as:

Predict / True	True 'spam'	True 'not spam'
Predict 'spam'	0	100
Predict 'not spam'	10	0

- Instead of assigning to most likely classify, **minimize expected cost**:

$$E[C(\hat{y}_i = \text{spam})] = p(y_i = \text{spam} | x_i) C(\hat{y}_i = \text{spam}, y_i = \text{spam}) + p(y_i = \text{not spam} | x_i) C(\hat{y}_i = \text{spam}, y_i = \text{not spam})$$

- Even if $p(\text{spam} | x_i) > p(\text{not spam} | x_i)$,

– Might still classify as “not spam”,

if $E[C(\hat{y}_i = \text{spam})] > E[C(\hat{y}_i = \text{not spam})]$.

"cost of predicting spam
when e-mail is not spam"

Decision Theory and Darts

- Post on decision theory in “darts”:
 - <http://www.datagenetics.com/blog/january12012/index.html>
- If you are very accurate, aim for the high-scoring regions.
- If you are very inaccurate, aim for the middle.
- Decision theory gives you the best strategy for other accuracies.

Summary

- **No free lunch theorem**: there is no “best” ML model.
- **Joint probability**: probability of A and B happening.
- **Conditional probability**: probability of A if we know B happened.
- **Generative classifiers**: build a probability of seeing the features.
 - Naïve Bayes uses **conditional independence** assumptions to make estimation practical.
- **Decision theory** allows us to consider costs of predictions.
- Next time:
 - A “best” machine learning model as ‘n’ goes to ∞ .



- All the remaining slides are “bonus”.
- We may go through them briefly, if time permits.

Generative Classifiers

- But does it need to know language to model $p(x_i | y_i)$???
- To fit generative models, usually make BIG assumptions:
 - Gaussian discriminant analysis (GDA):
 - Assume that $p(x_i | y_i)$ follows a multivariate normal distribution.
 - Naïve Bayes (NB):
 - Assume that each variables in x_i is independent of the others in x_i given y_i .

Bonus Slide: Avoiding Underflow

- During the prediction, the **probability can underflow**:

$$p(y=c | x_i) \propto \prod_{j=1}^d [p(x_{ij} | y=c)] p(y=c)$$

→ All these are < 1 so the product gets very small!

- Standard fix is to (equivalently) maximize the logarithm of the probability:

Remember that $\log(ab) = \log(a) + \log(b)$ so $\log(\prod a_i) = \sum \log(a_i)$

Since \log is monotonic the 'c' maximizing $p(y=c | x_i)$ also maximizes

$$\log p(y_i=c | x_i) = \sum_{j=1}^d [p(x_{ij} | y=c)] + p(y=c) + \underbrace{\text{constant}}_{\text{which is the same for all 'c'}}$$

Bonus Slide: $p(x_i)$ under naïve Bayes

- **Generative models** don't need $p(x_i)$ to make decisions.
- However, it's **easy to calculate** under the naïve Bayes assumption:

$$p(x_i) = \sum_{c=1}^K p(x_i, y=c) \quad (\text{marginalization rule})$$

$$= \sum_{c=1}^K p(x_i | y=c) p(y=c) \quad (\text{product rule})$$

$$= \sum_{c=1}^K \left[\prod_{j=1}^d p(x_{ij} | y=c) \right] p(y=c) \quad (\text{naïve Bayes assumption})$$

These are the quantities
we compute during training.

Bonus Slide: Less-Naïve Bayes

- Given features $\{x_1, x_2, x_3, \dots, x_d\}$, naïve Bayes approximates $p(y|x)$ as:

$$\begin{aligned} p(y | x_1, x_2, \dots, x_d) &\propto p(y) p(x_1, x_2, \dots, x_d | y) \quad \text{product rule applied repeatedly} \\ &= p(y) p(x_1 | y) p(x_2 | x_1, y) p(x_3 | x_2, x_1, y) \dots p(x_d | x_1, x_2, \dots, x_{d-1}, y) \\ &\approx p(y) p(x_1 | y) p(x_2 | y) p(x_3 | y) \dots p(x_d | y) \quad (\text{naïve Bayes assumption}) \end{aligned}$$

- The assumption is very strong, and there are “less naïve” versions:
 - Assume independence of all variables except up to ‘k’ largest ‘j’ where $j < i$.

- E.g., naïve Bayes has $k=0$ and with $k=2$ we would have:

$$\approx p(y) p(x_1 | y) p(x_2 | x_1, y) p(x_3 | x_2, x_1, y) p(x_4 | x_3, x_2, y) \dots p(x_d | x_{d-2}, x_{d-1}, y)$$

- Fewer independence assumptions so more flexible, but hard to estimate for large ‘k’.
 - Another practical variation is “tree-augmented” naïve Bayes.