

# CPSC 340: Machine Learning and Data Mining

Non-Linear Regression

# Admin

- **Assignment 1** grades are out.
- **Assignment 2** is due Sunday.
  - Extra office hours added on Saturday 12-2pm (see office hours calendar)
- **Assignment 3** will be out by early next week.
  - Will be due before the break
- **Midterm** is after the break (March 1 in class)
- Tutorial next week: practice problems for hw3

# Last Time: Linear Regression

- We discussed **linear models**:


$$y_i = w_1 x_{i1} + w_2 x_{i2} + \dots + w_d x_{id} \\ = \sum_{j=1}^d w_j x_{ij} = w^T x_i$$

- “Multiply feature  $x_{ij}$  by weight  $w_j$ , add them to get  $y_i$ ”.
- We discussed **squared error** function:

$$f(w) = \frac{1}{2} \sum_{i=1}^n (w^T x_i - y_i)^2$$

Predicted value  $\leftarrow$   $w^T x_i$        $\rightarrow$  True value  $y_i$

## The Social Cortex



DATA: THE SOCIAL BRAIN HYPOTHESIS, DUNBAR 1998

To predict on test case  $\hat{x}_i$   
use  $\hat{y}_i = w^T \hat{x}_i$

# Last Time: Supervised Learning Notation

- We're treating 'w', 'y', and each  $x_i$  as **column-vectors**:

$$w = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_d \end{bmatrix}_{d \times 1}$$

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}_{n \times 1}$$


$$x_i = \begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{id} \end{bmatrix}_{d \times 1}$$

- So feature matrix 'X' actually has  $x_i$  **transposed as rows**:

$$X = \begin{bmatrix} \text{---} x_1^T \text{---} \\ \text{---} x_2^T \text{---} \\ \vdots \\ \text{---} x_n^T \text{---} \end{bmatrix}$$


# Why don't we have a y-intercept?

- Last time: Linear models with **no y-intercept**.
  - Linear model is  $y_i = w^T x_i$  instead of  $y_i = w^T x_i + w_0$  with y-intercept  $w_0$ .
  - So if  $x_i = 0$  then we **must predict  $y_i = 0$** .



# Why don't we have a y-intercept?

- Last time: Linear models with **no y-intercept**.
  - Linear model is  $y_i = w^T x_i$  instead of  $y_i = w^T x_i + w_0$  with y-intercept  $w_0$ .
  - So if  $x_i = 0$  then we **must predict  $y_i = 0$** .



# Adding a Bias Variable

- Simple trick to add a y-intercept (“bias”) variable:
  - Make a new matrix “Z” with an **extra feature that is always “1”**.

$$X = \begin{bmatrix} 0.1 & 0.3 \\ 0.5 & -0.6 \\ 0.2 & 0.4 \end{bmatrix}$$

$$Z = \begin{bmatrix} 1 & 0.1 & 0.3 \\ 1 & 0.5 & -0.6 \\ 1 & 0.2 & 0.4 \end{bmatrix}$$

*(Note: A red bracket and 'X' mark are drawn under the last two columns of Z, indicating they are the original features from X.)*

- Now **use “Z” as features** to get a model with a **non-zero y-intercept**:

$$\begin{aligned} y_i &= w_0 z_{i0} + w_1 z_{i1} + w_2 z_{i2} \\ &\quad \quad \quad \downarrow \text{"1"} \quad \quad \downarrow x_{i1} \quad \quad \downarrow x_{i2} \\ &= w_0 + w_1 x_{i1} + w_2 x_{i2} \end{aligned}$$

- So we can have a **non-zero y-intercept by changing features**.

# Linear Least Squares

Prediction:

$$y = X * w$$

Why?

$$y_i = w^T x_i \quad \text{so}$$

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} w^T x_1 \\ w^T x_2 \\ \vdots \\ w^T x_n \end{bmatrix} = \begin{bmatrix} x_1^T w \\ x_2^T w \\ \vdots \\ x_n^T w \end{bmatrix} = \underbrace{\begin{bmatrix} \text{---} x_1^T \text{---} \\ \text{---} x_2^T \text{---} \\ \vdots \\ \text{---} x_n^T \text{---} \end{bmatrix}}_X \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_n \end{bmatrix} = Xw$$



# Linear Least Squares

- We can rewrite the objective function as a norm

Want ' $w$ ' that minimizes

$$f(w) = \frac{1}{2} \sum_{i=1}^n \underbrace{(w^T x_i - y_i)}_{r_i}^2 = \frac{1}{2} \sum_{i=1}^n r_i^2 = \frac{1}{2} r^T r = \frac{1}{2} \|r\|_2^2 = \boxed{\frac{1}{2} \|Xw - y\|_2^2}$$

Define "residual"  $r_i$  as

Signed error on example ' $i$ ':

$$r_i = w^T x_i - y_i$$

$$r = \begin{bmatrix} r_1 \\ r_2 \\ \vdots \\ r_n \end{bmatrix} = \begin{bmatrix} w^T x_1 - y_1 \\ w^T x_2 - y_2 \\ \vdots \\ w^T x_n - y_n \end{bmatrix} = \underbrace{\begin{bmatrix} w^T x_1 \\ w^T x_2 \\ \vdots \\ w^T x_n \end{bmatrix}}_{Xw} - \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \boxed{Xw - y}$$

# Linear Least Squares

Want 'w' that minimizes

$$f(w) = \frac{1}{2} \sum_{i=1}^n (w^T x_i - y_i)^2 =$$

$$\boxed{\frac{1}{2} \|Xw - y\|^2}$$

Let's expand  
then compute  
gradient.

$$= \frac{1}{2} (Xw - y)^T (Xw - y)$$

$$= \frac{1}{2} ((Xw)^T - y^T) (Xw - y)$$

$$= \frac{1}{2} (w^T X^T - y^T) (Xw - y)$$

$$= \frac{1}{2} (w^T X^T (Xw - y) - y^T (Xw - y))$$

$$= \frac{1}{2} (w^T X^T Xw - w^T X^T y - y^T Xw + y^T y)$$

$$= \frac{1}{2} w^T X^T Xw - w^T X^T y + \frac{1}{2} y^T y$$

A good way to check  
your steps: make sure  
that dimensions all make sense.

# Linear Least Squares

See notes on linear and quadratic derivatives for details.

Training:  $w = \text{solve}(X.T @ X, X.T @ y)$

Why?

Want 'w' that minimizes

$$f(w) = \frac{1}{2} \sum_{i=1}^n (w^T x_i - y_i)^2 = \frac{1}{2} \|Xw - y\|_2^2 = \frac{1}{2} w^T X^T X w - w^T X^T y + \frac{1}{2} y^T y$$

a quadratic function (in matrix notation)

What are the gradients of these terms?

$$\nabla f(w) = X^T X w - X^T y + 0$$

So at a minimizer where  $\nabla f(w) = 0$

we have:  $X^T X w = X^T y$  "normal equations"

Some matrix Some vector

This is a linear system  $Aw = b$

for some matrix 'A' and vector 'b'

Cheat sheet:  $\nabla_w [c] = 0$

$$\nabla_w [w^T b] = b$$

$$\nabla_w [\frac{1}{2} w^T A w] = A w \quad \text{for symmetric } A.$$

This is like saying  $\frac{d}{dw} [\alpha w] = \alpha$

# The Punch Line

$$\min_w \frac{1}{2} \|Xw - y\|_2^2$$



$$w = \text{solve}(X.T @ X, \quad X.T @ y)$$

*Note that  
f(w) is a "convex" function  
so solving  $\nabla f(w) = 0$  gives minimum*

# Incorrect Solutions to Least Squares Problem

The least squares objective is  $f(w) = \frac{1}{2} \|Xw - y\|^2$

The minimizers of this objective are solutions to the linear system:

$$X^T X w = X^T y$$

The following are not the solutions to the least squares problem:

$$w = (X^T X)^{-1} (X^T y) \quad (\text{only true if } \underline{X^T X \text{ is invertible}})$$

$$w X^T X = X^T y \quad (\text{matrix multiplication is } \underline{\text{not}} \text{ commutative, dimensions don't even match})$$

$$w = \frac{X^T y}{X^T X} \quad (\text{you } \underline{\text{cannot divide by a matrix}})$$

# Least Squares Issues

- Issues with least squares model:
  - Solution might **not be unique**.
  - It is **sensitive to outliers**.
  - It always **uses all features**.
  - Data can be so big we **can't store  $X^T X$** .
  - It might **predict outside range** of  $y_i$  values.
  - It assumes a **linear relationship** between  $x_i$  and  $y_i$ .

$X$  is  $n \times d$   
so  $X^T$  is  $d \times n$   
and  $X^T X$  is  $d \times d$ .


Costs  $O(nd^2)$  to calculate:

- Each of the  $O(d^2)$  elements is an inner product between length ' $n$ ' vectors.


# Example: Non-Linear Progressions in Athletics

- Are top athletes going faster, higher, and farther?


100m PROGRESSION MEN AND WOMEN (mean of top ten)



HIGH JUMP PROGRESSION MEN AND WOMEN (mean of top ten)



SHOT PUT PROGRESSION MEN (7.26 kg) AND WOMEN (4 kg) (mean of top ten)




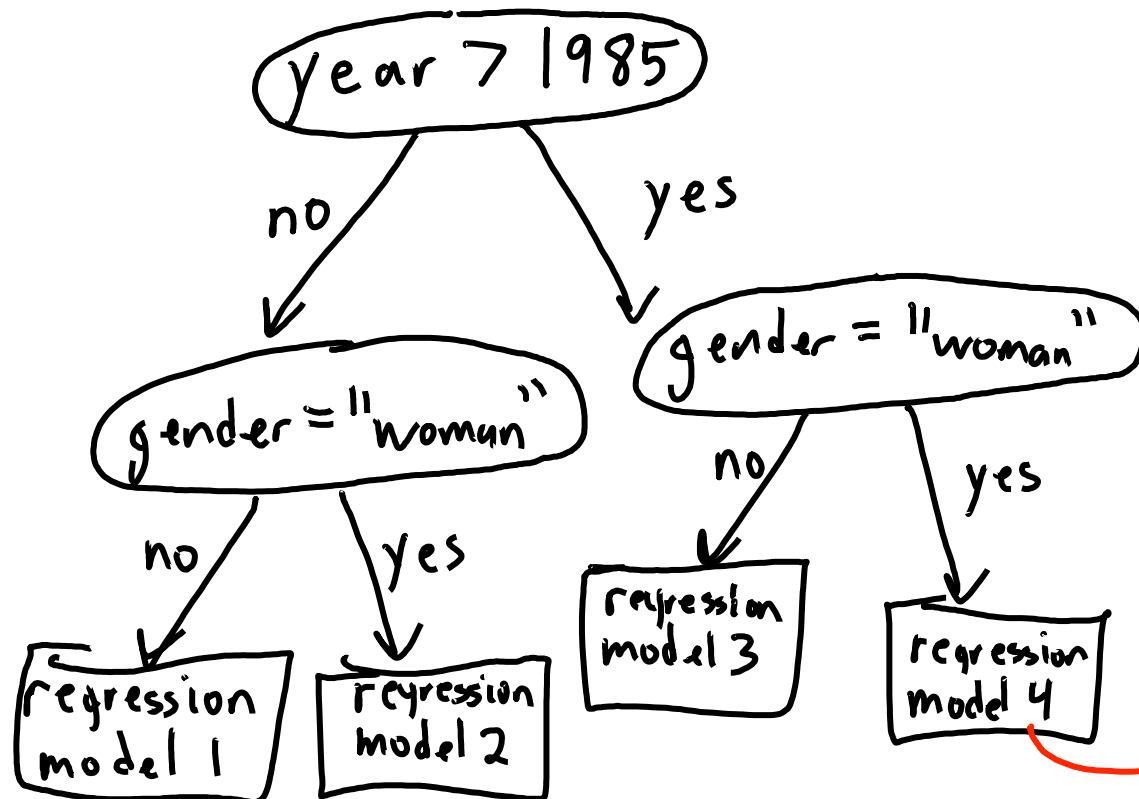
# Adapting Counting/Distance-Based Methods

- We can adapt our classification methods to perform regression:




# Adapting Counting/Distance-Based Methods

- We can adapt our classification methods to perform regression:
  - Regression tree: tree with mean value or linear regression at leaves.




# Adapting Counting/Distance-Based Methods

- We can adapt our classification methods to perform regression:
  - Regression tree: tree with mean value or linear regression at leaves.
  - Generative models: fit  $p(x_i | y_i)$  and  $p(y_i)$  with Gaussian or other model.




# Adapting Counting/Distance-Based Methods

- We can adapt our classification methods to perform regression:
  - Regression tree: tree with mean value or linear regression at leaves.
  - Generative models: fit  $p(x_i | y_i)$  and  $p(y_i)$  with Gaussian or other model.
  - Non-parametric models:
    - Mean  $y_i$  among k-nearest neighbours.



# Adapting Counting/Distance-Based Methods

- We can adapt our classification methods to perform regression:
  - Regression tree: tree with mean value or linear regression at leaves.
  - Generative models: fit  $p(x_i | y_i)$  and  $p(y_i)$  with Gaussian or other model.
  - Non-parametric models:
    - Mean  $y_i$  among k-nearest neighbours.
    - Could be weighted by distance.
      - Close points 'j' get more "weight"  $w_{ij}$ .



# Adapting Counting/Distance-Based Methods

- We can adapt our classification methods to perform regression:
  - Regression tree: tree with mean value or linear regression at leaves.
  - Generative models: fit  $p(x_i | y_i)$  and  $p(y_i)$  with Gaussian or other model.
  - Non-parametric models:
    - Mean  $y_i$  among  $k$ -nearest neighbours.
    - Could be weighted by distance.
  - Ensemble methods:
    - Can improve performance by averaging across regression models.

# Linear Least Squares for Quadratic Models

- Can we use **linear least squares** to fit a **quadratic model**?

$$y_i = w_0 + w_1 x_i + w_2 x_i^2$$

- You can do this by changing the features (**change of basis**):

$$X = \begin{bmatrix} 0.2 \\ -0.5 \\ 1 \\ 4 \end{bmatrix}$$

$$Z = \begin{bmatrix} 1 & 0.2 & (0.2)^2 \\ 1 & -0.5 & (-0.5)^2 \\ 1 & 1 & (1)^2 \\ 1 & 4 & (4)^2 \end{bmatrix}$$

$y$ -int
 $x$ 
 $x^2$


$$\begin{aligned} y_i &= w^T z_i \\ &= w_0 z_{i0} + w_1 z_{i1} + w_2 z_{i2} \\ &= w_0 + w_1 x_i + w_2 x_i^2 \end{aligned}$$

"solve  
linear  
system"

- Fitting with least squares:  $w = (Z^T Z)^{-1} (Z^T y)$
- It's a **linear function of  $w$** , but a **quadratic function of  $x_i$** .

To predict on new data  $\hat{X}$ , form  $\hat{Z}$  from  $\hat{X}$  and take  $y = \hat{Z}w$

# Linear Least Squares for Quadratic Models



# General Polynomial Basis



- We can have a polynomial of degree 'p' by using a basis:

$$Z = \begin{bmatrix} 1 & x_1 & (x_1)^2 & \dots & (x_1)^p \\ 1 & x_2 & (x_2)^2 & \dots & (x_2)^p \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & (x_n)^2 & \dots & (x_n)^p \end{bmatrix}$$

- There are polynomial basis functions that are numerically nicer:
  - E.g., Lagrange polynomials (see CPSC 303)

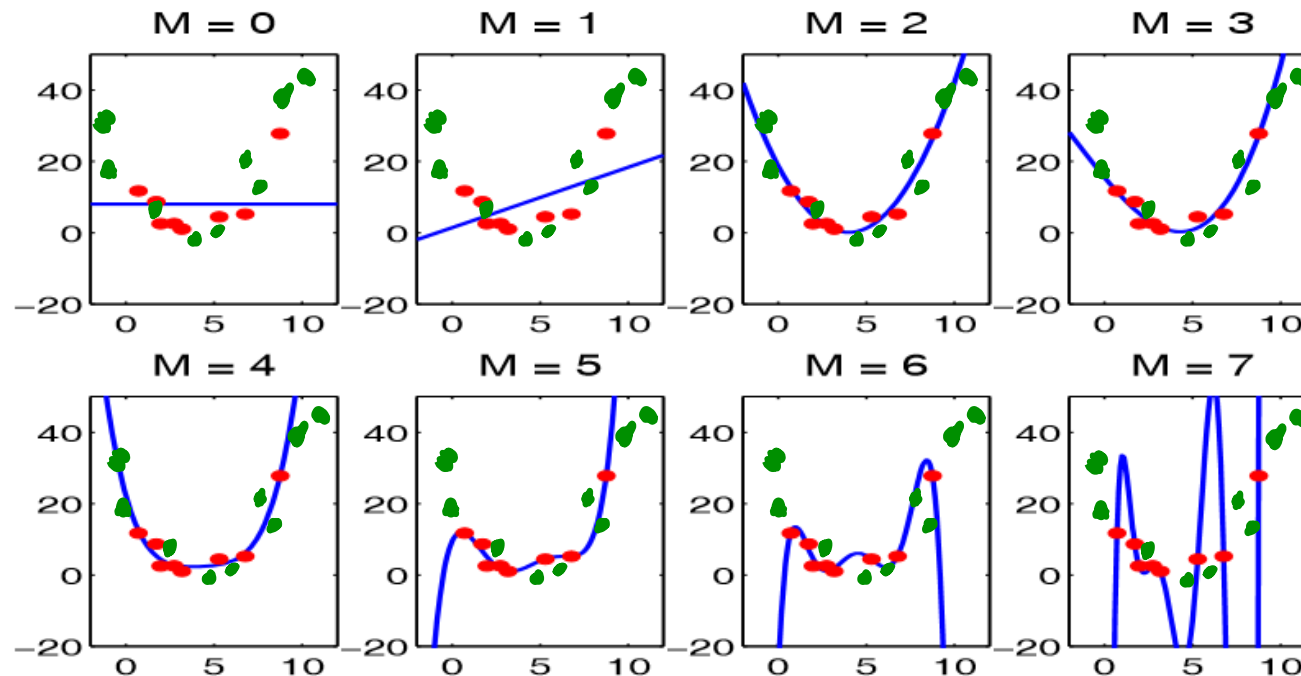


# General Polynomial Basis



# Degree of Polynomial and Fundamental Trade-Off

- As the polynomial degree increases, the **training error** goes down.



- But training error becomes worse approximation **test error**.
- Usual approach to **selecting degree**: **validation** or **cross-validation**.

# Summary

- Y-intercept can be modeled by using a column of 1s.
- Linear least squares solution is given by normal equations:
  - Solve  $(X^T X)w = X^T y$ .
- Tree/generative/non-parametric/ensemble methods for regression.
- Change of basis allows linear models to model non-linear data
- Next time:
  - A general method for avoiding overfitting.