

CPSC 340: Machine Learning and Data Mining

Gradient Descent

Admin

- **Assignment 2:**
 - Today is the last day to turn it in (with all late days).
 - We will release solutions tomorrow.
 - We won't make them available to other students
 - unless I hear from the person who wants to opt out
- **Assignment 3:**
 - Due in 11 days
- **Volunteers wanted: marking for Math Challengers competition**
 - This is a grade 8/9 math competition that I help organize (picking the questions)
 - Saturday Feb 18th 9:30am-2pm, at UBC
 - If interested, email mathchallengersvolunteers@gmail.com and CC me
 - (there are no incentives or anything, just if you actually want to do it)

Admin (continued)

- Grading disputes
 - If you perceive a problem with your homework grade, you have **1 week** to contact us via GitHub
 - Make sure to tag @cpsc340/staff or the particular graders listen on Piazza
 - Since this is a new policy, you have 1 week from today for hw1
 - More detailed info added to General Homework Instructions
- hw3 partnerships
 - Sorry, I'm going to stick to my guns here
 - If you're worried about this happening again, you can open all future issues now (for hw4, hw5, hw6)
 - You can still of course discuss the assignment, but will need to write separate solutions and code
 - I feel bad about this but need to stand up for myself at some point

Last Time: RBFs and Regularization

- We discussed **radial basis functions**:

- Basis functions that **depend on distances to training** points:

$$\begin{aligned} y_i &= w_1 \exp\left(-\frac{\|x_i - x_1\|^2}{2\sigma^2}\right) + w_2 \exp\left(-\frac{\|x_i - x_2\|^2}{2\sigma^2}\right) + \dots + w_n \exp\left(-\frac{\|x_i - x_n\|^2}{2\sigma^2}\right) \\ &= \sum_{j=1}^n w_j \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right) \end{aligned}$$

- Flexible bases that can **model any continuous function**.

- We also discussed **regularization**:

- Adding a **penalty on the model complexity**:

$$f(w) = \frac{1}{2} \|Xw - y\|^2 + \frac{\lambda}{2} \|w\|^2$$

- Best parameter lambda almost always leads to improved test error.
 - L2-regularized least squares is also known as “ridge regression”.

Features with Different Scales

- Consider features with different scales:

Egg (#)	Milk (mL)	Fish (g)	Pasta (cups)
0	250	0	1
1	250	200	1
0	0	0	0.5
2	250	150	0

- Should we convert to some standard 'unit'?
 - It **doesn't matter for least squares**:
 - $w_j \cdot (100 \text{ mL})$ gives the same model as $w_j \cdot (0.1 \text{ L})$
 - w_j will just be 1000 times smaller.
 - It also **doesn't matter for decision trees or naïve Bayes**.

Features with Different Scales

- Consider features with different scales:

Egg (#)	Milk (mL)	Fish (g)	Pasta (cups)
0	250	0	1
1	250	200	1
0	0	0	0.5
2	250	150	0

- Should we convert to some standard ‘unit’?
 - It **matters for k-nearest neighbours**:
 - KNN will focus on large values more than small values.
 - It **matters for regularized least squares**:
 - Penalization $|w_j|$ means different things if features ‘j’ are on different scales.

Standardizing Features

$$X = \begin{bmatrix} \vdots \end{bmatrix}$$

average of column 'j'

- It is common to **standardize features**:

- For each feature:

1. Compute mean and standard deviation:

$$\mu_j = \frac{1}{n} \sum_{i=1}^n x_{ij} \quad \sigma_j = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_{ij} - \mu_j)^2}$$

2. Subtract mean and divide by standard deviation:

Replace x_{ij} with $\frac{x_{ij} - \mu_j}{\sigma_j}$

- Means that **change in 'w_j' have similar effect** for any feature 'j'.
- Should we **regularize the bias**?
 - No! The y-intercept can be anywhere, why encourage it to be close to zero?
 - Yes! Regularizing all variables makes solution unique and it easier to compute 'w'.
 - Compromise: regularize the bias by a smaller amount than other variables?
 - I tried digging into the sklearn Ridge Regression code and it looks like "no".

Standardizing Target

- In regression, we sometimes **standardize the targets y_i** .
 - Puts targets on the same standard scale as standardized features:

$$\text{Replace } y_i \text{ with } \frac{y_i - \mu_y}{\sigma_y}$$

- With standardized target, setting $w = 0$ **predicts average y_i** :
 - High regularization makes us predict closer to the average value.
- Other common transformations of y_i are logarithm/exponent:

$$\text{Use } \log(y_i) \text{ or } \exp(\gamma y_i)$$

- Makes sense for geometric/exponential processes.

Ridge Regression Calculation

Objective: $f(w) = \frac{1}{2} \|Xw - y\|^2 + \frac{\lambda}{2} w^T w$

Gradient: $\nabla f(w) = X^T X w - X^T y + \lambda w$

Set $\nabla f(w) = 0$: $X^T X w + \lambda w = X^T y$

or
 $X^T X w + \lambda I w = X^T y$

or
 $(X^T X + \lambda I) w = X^T y$

Pre-multiply by $(X^T X + \lambda I)^{-1}$ which always exists:

$$\underbrace{(X^T X + \lambda I)^{-1} (X^T X + \lambda I)}_I w = (X^T X + \lambda I)^{-1} X^T y$$

Python:


```
w = solve(X.T@X+lam*np.eye(d), X.T*y)
```

so $w = (X^T X + \lambda I)^{-1} X^T y$ 9

Least Squares with Outliers

- Consider least squares problem with **outliers**:


$x \leftarrow$ "outlier" that doesn't follow trend



Least Squares with Outliers

- Consider least squares problem with **outliers**:


$x \leftarrow$ "outlier" that doesn't follow trend



- Least squares is very sensitive to outliers.

Least Squares with Outliers


- Squaring error shrinks small errors, and **magnifies large errors**:



- Outliers (large error) influence 'w' much more than other points.

Least Squares with Outliers

- Squaring error shrinks small errors, and **magnifies large errors**:



- Outliers (large error) influence 'w' much more than other points.
 - Good if outlier means 'plane crashes', bad if it means 'data entry error'.

Robust Regression

- Robust regression objectives put less focus large errors (outliers).
- For example, use absolute error instead of squared error:

$$f(w) = \sum_{i=1}^n |w^T x_i - y_i|$$

- Now decreasing 'small' and 'large' errors is equally important.
- Instead of minimizing L2-norm, minimizes L1-norm of residuals:

Least squares:

$$f(w) = \frac{1}{2} \|Xw - y\|^2$$


Least absolute error:

$$f(w) = \|Xw - y\|_1$$

$n \times d \quad d \times 1 \quad n \times 1$


Least Squares with Outliers

- Least squares is very sensitive to outliers.




Least Squares with Outliers

- Absolute error is more robust to outliers:



Regression with the L1-Norm

- Unfortunately, minimizing the absolute error is harder:
 - We can't take the gradient at zero.




- Generally, harder to minimize non-smooth than smooth functions.
- Could solve as 'linear program', but harder than 'linear system'.

Smooth Approximations to the L1-Norm

- There are **differentiable approximations to absolute value**.
- For example, the **Huber loss**:

$$f(w) = \sum_{i=1}^n h(w^T x_i - y_i)$$

$$h(r_i) = \begin{cases} \frac{1}{2} r_i^2 & \text{for } |r_i| \leq \epsilon \\ \epsilon(|r_i| - \frac{1}{2}\epsilon) & \text{otherwise} \end{cases}$$




- Setting $\nabla f(x) = 0$ does **not give a linear system**.
- But we can minimize 'f' using **gradient descent**:
 - Algorithm for finding local minimum of a differentiable function.

Gradient Descent for Finding a Local Minimum

- Gradient descent is an iterative optimization algorithm:
 - It starts with a “guess” w^0 .
 - It uses w^0 to generate a better guess w^1 .
 - It uses w^1 to generate a better guess w^2 .
 - It uses w^2 to generate a better guess w^3 .
 - ...
 - The limit of w^t as ‘t’ goes to ∞ has $\nabla f(w^t) = 0$.


Gradient Descent for Finding a Local Minimum

- Gradient descent is based on a simple observation:
 - Give parameters 'w', the direction of largest decrease is $-\nabla f(w)$.




Gradient Descent for Finding a Local Minimum

- Gradient descent is based on a simple observation:
 - Give parameters 'w', the direction of largest decrease is $-\nabla f(w)$.




Gradient Descent for Finding a Local Minimum

- Gradient descent is based on a simple observation:
 - Give parameters 'w', the direction of largest decrease is $-\nabla f(w)$.




Gradient Descent for Finding a Local Minimum

- Gradient descent is based on a simple observation:
 - Give parameters 'w', the direction of largest decrease is $-\nabla f(w)$.



Gradient Descent for Finding a Local Minimum

- Gradient descent is based on a simple observation:
 - Give parameters 'w', the direction of largest decrease is $-\nabla f(w)$.



Now the slope $\nabla f(w^4)$ is positive
so we move in the negative direction.

Gradient Descent for Finding a Local Minimum

- Gradient descent is an **iterative optimization** algorithm:
 - It starts with a “guess” w^0 .
 - Generate new guess by moving in the negative gradient direction:

$$w^1 = w^0 - \alpha^0 \nabla f(w^0)$$

- This decreases f if the step size α is small enough
- Repeat to successively refine the guess


$$w^{t+1} = w^t - \alpha^t \nabla f(w^t) \quad \text{for } t = 1, 2, 3, \dots$$

- Stop if not making progress or

$$\|\nabla f(w^t)\| \leq \delta$$

Some small scalar.
Approximate local minimum




Gradient Descent in 2D



- Under weak conditions, algorithm converges to a local minimum.

Convex Functions

- Is finding a **local minimum good enough**?
 - For least squares and Huber loss this is enough: they are **convex functions**.



- A function is **convex** if the **area above the function is a convex set**.
 - All values between any two points above function stay above function.


Convex Functions

- All local minima of convex functions are also global minima.

Proof by contradiction:

Consider a local minimum

If this is not global minimum,
there must a smaller value.



But this
contradicts that
we are at a
local minimum.

By convexity we can move along line to global minimum and decrease objective.

- Gradient descent finds a global minimum on convex functions.
- Next time: how do we know if a function is convex?

Gradient Descent

- Least squares via normal equations vs. gradient descent:

- Normal equations cost $O(nd^2 + d^3)$.

Forming $X^T X$ costs $O(nd^2)$ and solving a $d \times d$ linear system costs $O(d^3)$

- Gradient descent costs $O(ndt)$ to run for 't' iterations.

Computing $\nabla f(w) = X^T X w - X^T y$ only costs $O(nd)$.

→ We can use $X^T X w = X^T (X w)$ which is just two $n \times d$ matrix multiplications.

- Gradient descent can be faster when 'd' is very large:
- Improving on gradient descent: Nesterov and Newton method.
 - For L2-regularized least squares, there is also “conjugate” gradient.


Motivation for Considering Worst Case



'Brittle' Regression

- What if you really care about **getting the outliers right?**
 - You want **best performance on worst training example.**
 - For example, if in worst case the plane can crash.
- In this case you can use something like the infinity-norm:

$$f(w) = \|Xw - y\|_{\infty} \quad \text{where} \quad \|r\|_{\infty} = \max_i \{ |r_i| \}$$



- Very sensitive to outliers (brittle), but worst case will be better.

Log-Sum-Exp Function

- As with the L_1 -norm, the L_∞ norm is convex but non-smooth:
 - We can fit it with gradient descent using a smooth approximation.
- Log-sum-exp function is a smooth approximation to the max function:

$$\max_i \{z_i\} \approx \log\left(\sum_i \exp(z_i)\right)$$

- Intuition: largest element is magnified exponentially.
 - Smaller elements become negligible in comparison.
 - Recall that $\log(\exp(z))=z$.
- Notation alert: by “log” I always mean the natural logarithm.

Summary

- Robust regression using L1-norm/Huber is less sensitive to outliers.
- Gradient descent finds local minimum of differentiable function.
- Convex functions do not have non-global local minima.
- Log-Sum-Exp function: smooth approximation to maximum.
- Next time:
 - Finding ‘important’ e-mails, and beating naïve Bayes on spam filtering.