

CPSC 340: Machine Learning and Data Mining

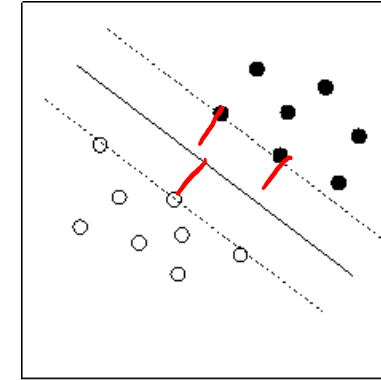
Kernel Methods

Admin

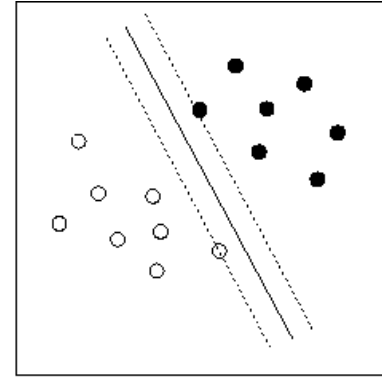
- **Assignment 3:**
 - Due Sunday evening
 - Solutions will be posted on Thursday
- **Assignments 1 and 2:**
 - You can now see each other's work
- **Midterm March 1**
 - Past exams posted
 - Midterm covers Assignments 1-3 / lectures 1-16
 - Tutorials after break will cover practice exam questions
 - In class, 1pm-1:55pm, closed-book, 1 page double-sided “cheat sheet”.

Last Time: SVMs and Kernel Trick

- We discussed the **maximum margin** view of **SVMs**:
 - Yields an **L2-regularized hinge loss**.



(a) Larger margin



(b) Smaller margin

- We introduced the **kernel trick**:
 - Write model to only depend on **inner products between features vectors**.

$$\hat{y} = \hat{K} (K + \lambda I)^{-1} y$$

$t \times n$ matrix $\hat{Z}Z^T$ containing inner products between test examples and training examples. \rightarrow $n \times n$ matrix ZZ^T containing inner products between all training examples.

- So everything we need to know about z_i is summarized by the $z_i^T z_j$.
- If you have a **kernel function** $k(x_i, x_j)$ that computes $z_i^T z_j$, then you don't need to compute the basis z_i explicitly.

Polynomial Kernel with Higher Degrees

- Assume that I have 2 features and want to use the **degree-2 basis**:

$$z_i = [1 \quad \sqrt{2}x_{i1} \quad \sqrt{2}x_{i2} \quad x_{i1}^2 \quad \sqrt{2}x_{i1}x_{i2} \quad x_{i2}^2]^T$$

- I can compute **inner products** using:

$$\begin{aligned} (1 + x_i^T x_j)^2 &= 1 + 2x_i^T x_j + (x_i^T x_j)^2 \\ &= 1 + 2x_{i1}x_{j1} + 2x_{i2}x_{j2} + x_{i1}^2x_{j1}^2 + 2x_{i1}x_{i2}x_{j1}x_{j2} + x_{i2}^2x_{j2}^2 \end{aligned}$$

$$\begin{aligned} &= \underbrace{[1 \quad \sqrt{2}x_{i1} \quad \sqrt{2}x_{i2} \quad x_{i1}^2 \quad \sqrt{2}x_{i1}x_{i2} \quad x_{i2}^2]}_{z_i^T} \underbrace{\begin{bmatrix} 1 \\ \sqrt{2}x_{j1} \\ \sqrt{2}x_{j2} \\ x_{j1}^2 \\ \sqrt{2}x_{j1}x_{j2} \\ x_{j2}^2 \end{bmatrix}}_{z_j} \\ &= z_i^T z_j \end{aligned}$$

Polynomial Kernel with Higher Degrees

- To get all degree-4 “monomials” I can use:

$$z_i^T z_j = (x_i^T x_j)^4$$

Equivalent to using a z_i with weighted versions of $x_{i1}^4, x_{i1}^3 x_{i2}, x_{i1}^2 x_{i2}^2, x_{i1} x_{i2}^3, x_{i2}^4, \dots$

- To also get lower-order terms use $z_i^T z_j = (1 + x_i^T x_j)^4$
- The general degree- p **polynomial kernel** function:

$$k(x_i, x_j) = (1 + x_i^T x_j)^p$$

- Works for any number of features ‘ d ’.
- But cost of computing $z_i^T z_j$ is $O(d)$ instead of $O(d^p)$.
- Take-home message: I can take the dot products without constructing the feature vectors themselves.

Kernel Trick

- Using polynomial basis of degree 'p' with the kernel trick:

- Compute K and \hat{K} :

$$K_{ij} = (1 + x_i^T x_j)^p \quad \hat{K}_{ij} = (1 + \hat{x}_i^T x_j)^p$$

\hat{x}_i ← test example
 x_j ← train example

- Make predictions using:

$$\hat{y} = \hat{K} (K + \lambda I)^{-1} y$$

\hat{y} (tx1) \hat{K} (txn) $(K + \lambda I)^{-1}$ (nxn) y (nx1)

\nwarrow To form $K = Z Z^T$

- Training cost is only $O(n^2 d + n^3)$, despite using $O(d^p)$ features.

- Testing cost is only $O(ndt)$.

\nwarrow To invert nxn matrix
 \nwarrow To form $\hat{K} = \hat{Z} Z^T$

Linear Regression vs. Kernel Regression

Linear Regression

Training

1. Form basis Z from X
2. Compute $w = (Z^T Z + \lambda I)^{-1} (Z^T y)$

Testing

1. Form basis \hat{Z} from \hat{X}
2. Compute $\hat{y} = \hat{Z} w$

Kernel Regression

Training

1. Form inner products K from X .
2. Compute $v = (K + \lambda I)^{-1} y$

Testing:

1. Form inner products \hat{K} from X and \hat{X}
2. Compute $\hat{y} = \hat{K} v$

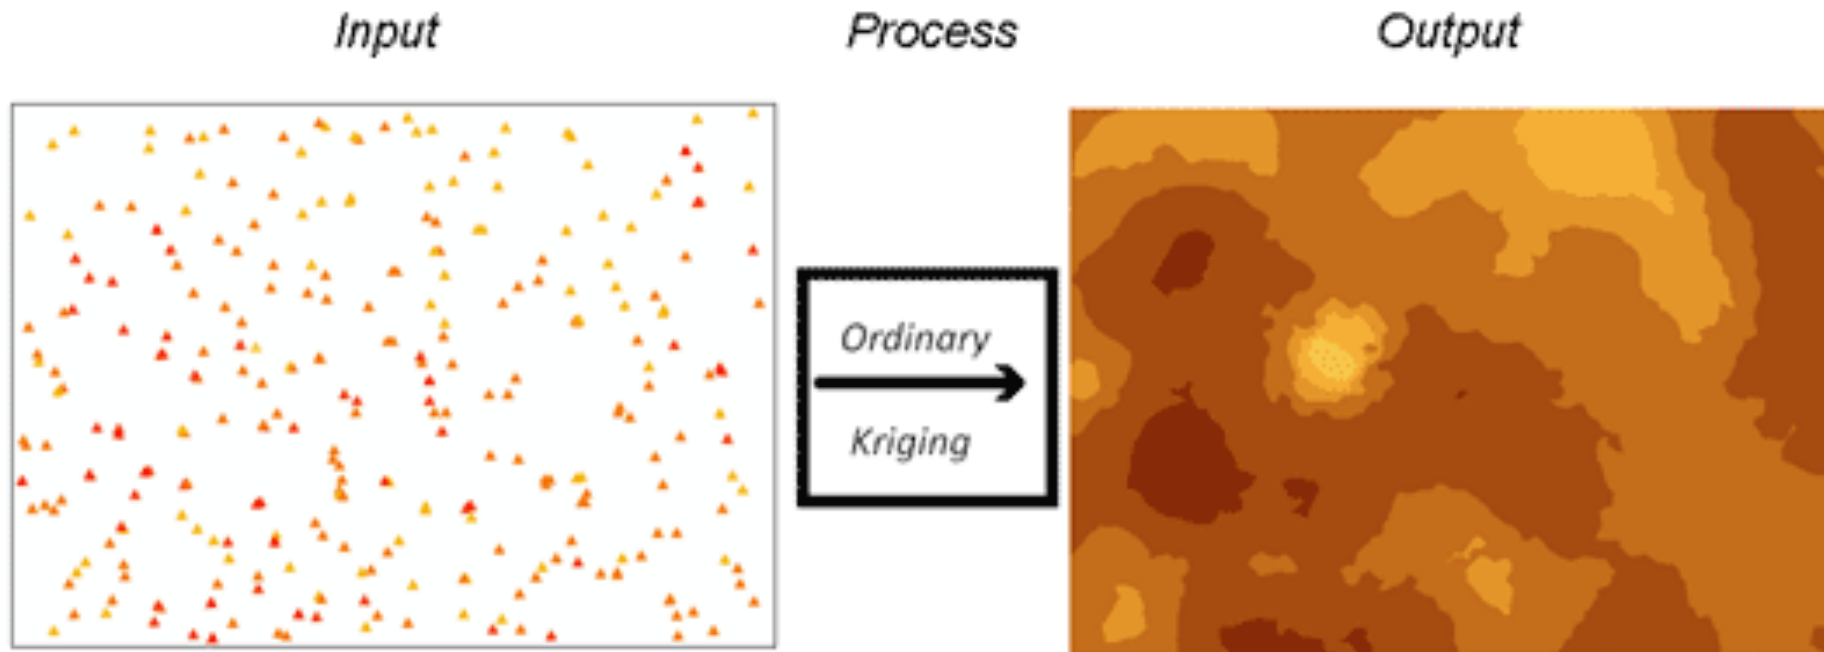
Non-parametric



Observation: this requires all training examples ☹

Motivation: Finding Gold

- Kernel methods first came from mining engineering ('Kriging'):
 - Mining company wants to find gold.
 - Drill holes, measure gold content.
 - Build a kernel regression model (typically use RBF kernels).



Gaussian-RBF Kernel

- Most common kernel is the **Gaussian RBF** kernel:

$$k(x_i, x_j) = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right)$$

- Same formula and behaviour as RBF basis, but not equivalent:
 - Before we used RBFs as a basis, now we're using them as inner-product.
- Basis z_i giving the **Gaussian RBF kernel is infinite-dimensional**.
 - Not much hope of doing this without the kernel trick...
- Kernel trick lets us **fit regression models without explicit features**:
 - We can interpret $k(x_i, x_j)$ as a “similarity” between objects x_i and x_j .
 - We **don't need z_i and z_j** if we can compute ‘similarity’ between objects.

Kernel trick for structured data

- Consider data that doesn't look like this:

$$X = \begin{bmatrix} 0.5377 & 0.3188 & 3.5784 \\ 1.8339 & -1.3077 & 2.7694 \\ -2.2588 & -0.4336 & -1.3499 \\ 0.8622 & 0.3426 & 3.0349 \end{bmatrix}, \quad y = \begin{bmatrix} +1 \\ -1 \\ -1 \\ +1 \end{bmatrix},$$

- But instead looks like this:

$$X = \begin{bmatrix} \text{Do you want to go for a drink sometime?} \\ \text{J'achète du pain tous les jours.} \\ \text{Fais ce que tu veux.} \\ \text{There are inner products between sentences?} \end{bmatrix}, y = \begin{bmatrix} +1 \\ -1 \\ -1 \\ +1 \end{bmatrix}.$$

- Instead of using features, can **define kernel between sentences**.
 - E,g, “string kernels”: weighted frequency of common subsequences.
- There are also “image kernels”, “graph kernels”, and so on...

Valid Kernels

- What kernel functions $k(x_i, x_j)$ can we use?
- Kernel 'k' must be an inner product in some space:
 - There must exist a mapping from x_i to some z_i such that $k(x_i, x_j) = z_i^T z_j$.
- It can be hard to show that a function satisfies this.
- But there are some simple rules for constructing valid kernels from other valid kernels (bonus slide).

Kernel Trick for Other Methods

- Besides **L2-regularized least squares**, when can we use kernels?

– **Methods based on Euclidean distances** between examples:

- Kernel k-nearest neighbours.
- Kernel clustering (k-means, DBSCAN, hierarchical).
- Kernel outlierness.
- Kernel “Amazon Product Recommendation”.
- Kernel non-parametric regression.

$$\|z_i - z_j\|^2 = z_i^T z_i - 2z_i^T z_j + z_j^T z_j = k(x_i, x_i) - 2k(x_i, x_j) + k(x_j, x_j)$$

– **L2-regularized linear models** (“representer theorem” -- see bonus slides):

- L2-regularized robust regression.
- L2-regularized logistic regression.
- L2-regularized support vector machines.

less obvious but true

With a particular implementation testing cost is reduced from $O(ndt)$ to $O(mdt)$ Number of support vectors.

Kernel trick continued

- Because of the support vectors, kernels are used with SVMs quite often, but much less so with logistic regression.

`sklearn.svm.SVC`

```
class sklearn.svm. SVC (C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, probability=False, tol=0.001, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape=None, random_state=None)
```

[\[source\]](#)

`sklearn.linear_model.LogisticRegression`

```
class sklearn.linear_model. LogisticRegression (penalty='l2', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None, solver='liblinear', max_iter=100, multi_class='ovr', verbose=0, warm_start=False, n_jobs=1)
```

[\[source\]](#)

RBF kernel vs RBF features

- Like the RBF features, the RBF kernel...
 - can learn any decision boundary given enough data
 - as a result it is prone to overfitting, so we need to use regularization
 - σ parameter controls smoothness: larger σ means smoother boundaries
 - This is called "gamma" in sklearn and it's $1/\sigma$
 - λ parameter controls regularization: larger λ means more regularization
 - This is called "C" in sklearn and it's $1/\lambda$
- The RBF features are finite-dimensional (N features)
- The RBF kernel corresponds to infinitely many features
- Both are non-parametric methods

Summary

- Kernels let us use similarity between objects, rather than features.
- The RBF kernel allows us to use infinitely many features in finite computational time.
- We'll spend the rest of today's class reviewing recent topics.

Bonus Slide: Features Corresponding to RBF Kernel

Gaussian-RBF Kernels

- The most common kernel is the **Gaussian-RBF** (or 'squared exponential') kernel,

$$k(x_i, x_j) = \exp\left(-\frac{\|x_i - x_j\|^2}{\sigma^2}\right).$$

- What function $\phi(x)$ would lead to this as the inner-product?
 - To simplify, assume $d = 1$ and $\sigma = 1$,

$$\begin{aligned}k(x_i, x_j) &= \exp(-x_i^2 + 2x_i x_j - x_j^2) \\ &= \exp(-x_i^2) \exp(2x_i x_j) \exp(-x_j^2),\end{aligned}$$

so we need $\phi(x_i) = \exp(-x_i^2)z_i$ where $z_i z_j = \exp(2x_i x_j)$.

- For this to work for *all* x_i and x_j , z_i must be infinite-dimensional.
- If we use that

$$\exp(2x_i x_j) = \sum_{k=0}^{\infty} \frac{2^k x_i^k x_j^k}{k!},$$

then we obtain

$$\phi(x_i) = \exp(-x_i^2) \left[1 \quad \sqrt{\frac{2}{1!}} x_i \quad \sqrt{\frac{2^2}{2!}} x_i^2 \quad \sqrt{\frac{2^3}{3!}} x_i^3 \quad \cdots \right].$$

Bonus Slide: Designing Valid Kernel Functions

Constructing Valid Kernels

- If $k_1(x_i, x_j)$ and $k_2(x_i, x_j)$ are valid kernels, then the following are valid kernels:
 - $k_1(\phi(x_i), \phi(x_j))$.
 - $\alpha k_1(x_i, x_j) + \beta k_2(x_i, x_j)$ for $\alpha \geq 0$ and $\beta \geq 0$.
 - $k_1(x_i, x_j)k_2(x_i, x_j)$.
 - $\phi(x_i)k_1(x_i, x_j)\phi(x_j)$.
 - $\exp(k_1(x_i, x_j))$.
- Example: Gaussian-RBF kernel:

$$\begin{aligned} k(x_i, x_j) &= \exp\left(-\frac{\|x_i - x_j\|^2}{\sigma^2}\right) \\ &= \underbrace{\exp\left(-\frac{\|x_i\|^2}{\sigma^2}\right)}_{\phi(x_i)} \underbrace{\exp\left(\underbrace{\frac{2}{\sigma^2}}_{\alpha \geq 0} \underbrace{x_i^T x_j}_{\text{valid}}\right)}_{\exp(\text{valid})} \underbrace{\exp\left(-\frac{\|x_j\|^2}{\sigma^2}\right)}_{\phi(x_j)}. \end{aligned}$$

Bonus Slide: Kernels for Linear Model plus L2-Reg

Representer Theorem

- Consider linear model differentiable with losses f_i and L2-regularization,

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \sum_{i=1}^n f_i(w^T x_i) + \frac{\lambda}{2} \|w\|^2.$$

- Setting the gradient equal to zero we get

$$0 = \sum_{i=1}^n f'_i(w^T x_i) x_i + \lambda w.$$

- So any solution w^* can be written as a **linear combination of features x_i** ,

$$\begin{aligned} w^* &= -\frac{1}{\lambda} \sum_{i=1}^n f'_i((w^*)^T x_i) x_i = \sum_{i=1}^n z_i x_i \\ &= X^T z. \end{aligned}$$

- This is called a **representer theorem** (true under much more general conditions).

Bonus Slide: Kernels for Linear Model plus L2-Reg

Representer Theorem

- Using representer theorem we can use $w = X^T z$ in original problem,

$$\begin{aligned} & \operatorname{argmin}_{w \in \mathbb{R}^d} \sum_{i=1}^n f_i(w^T x_i) + \frac{\lambda}{2} \|w\|^2 \\ &= \operatorname{argmin}_{z \in \mathbb{R}^n} \sum_{i=1}^n f_i(\underbrace{z^T X x_i}_{x_i^T X^T z}) + \frac{\lambda}{2} \|X^T z\|^2 \end{aligned}$$

- Now defining $f(z) = \sum_{i=1}^n f_i(z_i)$ for a vector z we have

$$\begin{aligned} &= \operatorname{argmin}_{z \in \mathbb{R}^n} f(X X^T z) + \frac{\lambda}{2} z^T X X^T z \\ &= \operatorname{argmin}_{z \in \mathbb{R}^n} f(Kz) + \frac{\lambda}{2} z^T K z. \end{aligned}$$

- Similarly, at test time we can use the n variables z ,

$$\hat{X} w = \hat{X} X^T z = \hat{K} z.$$